

Institute of Information Systems

Distributed Systems Group (DSG)
UE Verteilte Systeme WS 2016/17(184.167)

WS 2016/17
Lab 2

Submission Deadline: 11.01.2017, 18:00

Contents

1	General Remarks	3
2	Submission Guide	4
2.1	Submission	4
2.2	Interviews	4
3	Installation/Registration of the Bouncy Castle Provider	5
4	Project Template	5
5	Lab Description	6
5.1	Stage 1 - Naming service and RMI (16 points)	6
5.1.1	Description	6
5.1.2	Overview	6
5.1.3	Nameserver	8
5.1.4	Chatserver Updates	10
5.2	Stage 2 - Secure channel (12 points)	12
5.2.1	Authentication Algorithm	12
5.2.2	Client Application Behavior	13
5.2.3	Chatserver Application Behavior	14
5.3	Stage 3 - Message Integrity (7 points)	15
6	Lab Port Policy	16
7	Regular expressions	16
8	Further Reading Suggestions	17

1 General Remarks

- We suggest reading (at least) the following materials before you start implementing:
 - Chapter 9 - Security from the book Distributed Systems: Principles and Paradigms (2nd edition)
 - Java Cryptography Architecture (JCA) Reference Guide¹: Tutorial about the Java Cryptography Architecture (JCA).
 - Java RMI Tutorial²: A short introduction into RMI.
 - JGuru RMI Tutorial³: A more detailed tutorial about RMI.
- **Lab 2 is a group assignment** - make sure to distribute the work evenly within your group, and also make sure that **all group members** fully understand the entire solution in detail. Each group member should be capable of describing the code, also the parts which they did not implement!
- Collaboration across group boundaries (i.e., exchanging code with other groups) is **NOT** allowed. Discussions with colleagues (e.g., in the forum) are allowed but the code has to be written alone.
- Don't use any other 3rd party library except the ones we provided for you (e.g., Bouncy Castle).
- Be sure to check the Hints & Tricky Parts⁴ section! For this assignment we have put lots of helpful code snippets there!
- Furthermore check the provided FAQ⁵, where we discuss a lot of known issues and problems from the last years!

¹<http://docs.oracle.com/javase/7/docs/technotes/guides/security/crypto/CryptoSpec.html>

²<http://docs.oracle.com/javase/tutorial/rmi/index.html>

³<http://www.dse.disco.unimib.it/ds/extra/rmiTutorial.pdf>

⁴<https://tuwel.tuwien.ac.at/mod/book/view.php?id=277009&chapterid=118>

⁵<https://tuwel.tuwien.ac.at/mod/book/view.php?id=277009&chapterid=103>

2 Submission Guide

2.1 Submission

- You must upload your solution using TUWEL before the submission deadline: **12.01.2017, 18:00** - please note that the deadline is hard! You are responsible for submitting your solution in time. If you do not submit, you won't get any points!
- Upload your solution as a **ZIP file**. Please submit **only the provided template and your classes**, the `build.xml` file and a `readme.txt` (no compiled class files, no third-party libraries - except the libraries already provided in the template, no svn/git metadata, no hidden files etc.).
- The purpose of `readme.txt` is to reflect about your solution. It should contain a short summary of the status of your code so that a tutor can get the information right before the mandatory interview (see below) and can give you some tips for the next assignment.
- Your submission must compile and run in our lab environment. Use and complete the project template provided in TUWEL.
- Test your solution extensively in our lab environment. It'll be worth the time.
- Please make sure that your upload was successful (i.e., you should be able to download your solution - as the tutors will do during the interview).

2.2 Interviews

- After the submission deadline, there will be a mandatory interview (Abgabegespräch). You must register for a time slot for the interviews using TUWEL.
- You can do the interview only if you submitted your solution before the deadline!
- The interview will take place in the DSLab PC room⁶. During the interview, you will be asked about the solution that you uploaded (i.e., **changes after the deadline will not be taken into account!**). In the interview you need to explain your code, design and architecture in detail.
- Remember that you can do the interview only once!

Important: Lab 2 consists of several (more or less independent) stages. If you want to get all points for this assignment, you will have to implement all of them. If you are satisfied with less, you can leave one or more of them unimplemented.

⁶<https://tuwel.tuwien.ac.at/mod/book/view.php?id=277009&chapterid=94>

3 Installation/Registration of the Bouncy Castle Provider

For the second and third part we will use the Bouncy Castle⁷ library as a provider for the Java Cryptography Extension (JCE) API, which is part of JCA. The Bouncy Castle provider (JDK 1.6 version) is already part of the provided template. Please stick to the provided version as this is the one used in our lab environment.

The provider is configured as part of your Java environment by adding an entry to the `java.security` properties file (found in `$JAVA_HOME/jre/lib/security/java.security`, where `$JAVA_HOME` is the location of your JDK distribution). You will find detailed instructions in the file, but basically it comes down to adding this line (but you may need to move all other providers one level of preference up):

```
security.provider.1 = org.bouncycastle.jce.provider.BouncyCastleProvider
```

Where you actually put the jar file is mostly up to you, but the best place to have it is in `$JAVA_HOME/jre/lib/ext` (which is the **jre folder within the JDK installation** - do not confuse with the JRE installation).

The installation of a custom provider is explained in the Java Cryptography Architecture (JCA) Reference Guide⁸ in detail.

Note: If you get "`java.lang.SecurityException: Unsupported keysize or algorithm parameters`" or "`java.security.InvalidKeyException: Illegal key size`" exception while using the Bouncy Castle library, then check this hint⁹.

4 Project Template

The `lib` directory contains the Bouncy Castle library. The `keys` directory contains all the keys required to test your implementation. The private keys used in the communication between client and chatserver are encrypted with following password for the respective participant:

```
alice.vienna.at: 12345  
bill.de: 23456  
chatserver: 12345
```

The secret key used in the client-to-client communication is not password protected at all.

Please check again that you modified several port constants in the `.properties` files according to your assigned port range from Lab 0.

⁷<http://www.bouncycastle.org/>

⁸<http://docs.oracle.com/javase/7/docs/technotes/guides/security/crypto/CryptoSpec.html#ProviderInstalling>

⁹<https://tuwel.tuwien.ac.at/mod/book/view.php?id=277009&chapterid=118>

5 Lab Description

Lab 2 is divided in the following parts:

- **Stage 1** (16 points): In this stage you will implement a simple distributed naming service that uses **Java Remote Method Invocation (RMI)**. The distributed naming service extends the mechanism for the registration and lookup of private user addresses from the first lab. For this, you will get acquainted with basic naming¹⁰, name server lookup, and RMI communication.
- **Stage 2** (12 points): In this stage you will create a security layer for the TCP communication between the client and the chatserver by implementing a secure channel and mutual authentication using **public-key cryptography**. To that end, you will use the **Java Cryptography Architecture (JCA)**. In our case, the secure channel will protect both parties against interception and fabrication of messages. Note that the common definition of a secure channel additionally implies a protection against modification. The global communication will remain vulnerable in this regard; however, in Stage 3, we address this issue for private communication between clients.
- **Stage 3** (7 points): In this stage you will use JCA to secure the communication against message modification. For the sake of simplicity, we will add this feature solely to the otherwise unsecured private communication between clients. By using a **Message Authentication Code (MAC)** and prepending it to each message, the receiver can check whether the message has been modified on the way through the channel.

5.1 Stage 1 - Naming service and RMI (16 points)

5.1.1 Description

In this stage you will learn:

- the basics of a simple distributed object technology (RMI)
- how to bind and lookup objects with a naming service
- how to implement callbacks with RMI.

5.1.2 Overview

In this stage you will extend the chatserver from lab 1 by adding a distributed and hierarchical naming service. Currently, the chatserver uses a primitive in-memory mechanism for storing and looking up private addresses of users. As the network grows, this approach would quickly turn the chatserver into a bottleneck. To allow scalability of the naming system, the naming service should be updated to use a *network* of nameservers. Figure 1 shows a simplified overview of the updated system architecture.

Similar to DNS, our nameservers span a distributed namespace as shown in Figure 1. The namespace network is hierarchically divided into a collection of domains. There is a single top-level domain, hosted by the *root nameserver*. Using nameservers, each domain can be divided into smaller subdomains. Each domain is managed by exactly one nameserver. Likewise, nameservers only host exactly one zone. A nameserver can communicate with the nameservers on the next lower level. In our system, nameservers never need to contact their parent nameservers, i.e., name resolution is only performed top-down.

In our scenario, a domain name is composed of a sequence of zones, where each zone consists of alphabetic characters (case insensitive). If a domain consists of multiple zones, these zones are separated by single dots ('.'). In our case, 'berlin.de' is a domain, as are 'rome.it' and 'at'. Saying 'berlin' is a *zone* in the namespace of 'de' is the same like saying 'berlin' is a subdomain of 'de'.

¹⁰The naming strategy that you will implement is presented in this assignment. For more information about naming in general, please consult Chapter 5 from the book Distributed Systems: Principles and Paradigms (2nd edition).

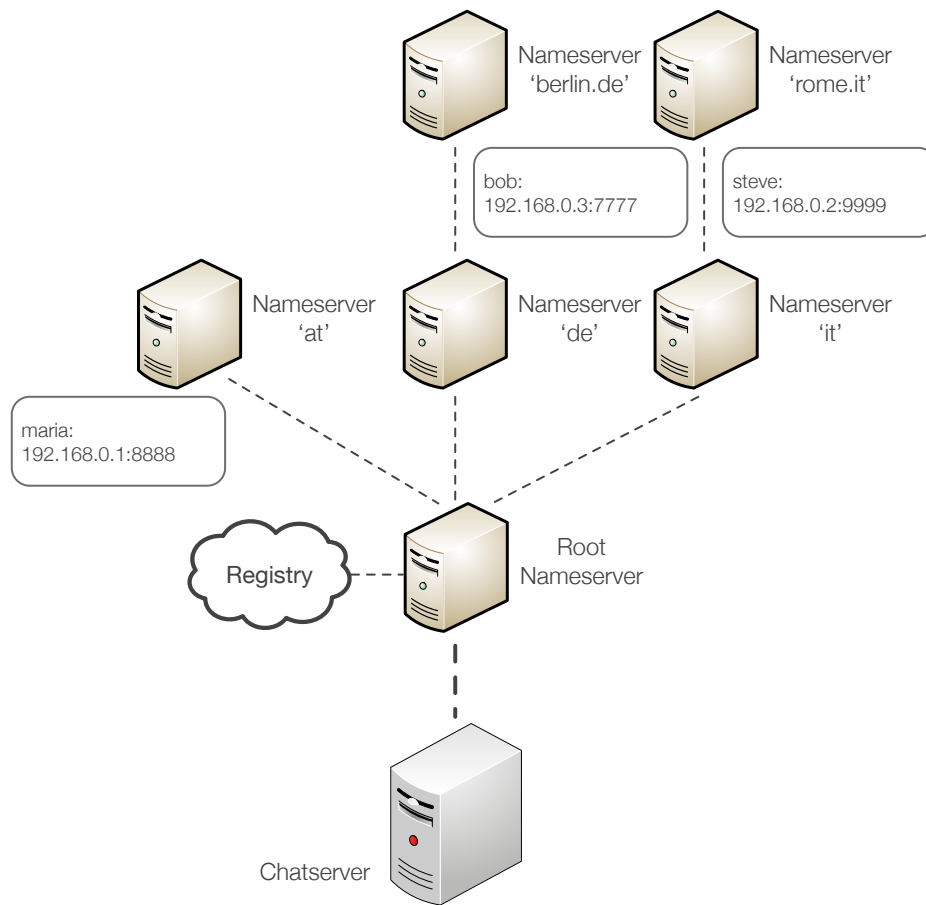


Figure 1: Naming Service: Overview

Registration of new Domains When a new subdomain should be added to the system, the new nameserver has to be registered with the respective parent name server first. Domain registrations are carried out in a recursive manner, illustrated in Figure 2. First, the new nameserver contacts the root nameserver, attaching the name of the desired domain to its request. In the example shown in Figure 2, the new nameserver tries to register the domain 'berlin.de'. The root nameserver then forwards the request to the next lower level respectively. This procedure continues until further name resolution is no longer possible. Accordingly, in the third step of the example, the request is sent to the nameserver hosting the 'de'-domain. At this point, there is only one subdomain left ('berlin'), which must be the requested zone. Therefore, the 'de'-nameserver can store the new server as a child responsible for the new zone 'berlin'.

Bootstrapping One well-known issue in such scenarios is the bootstrapping problem: To connect to a network, the respective participant needs to know the address of at least one other participant already connected to the network. In our case, to start the registration process, a new server first needs to find the root nameserver. To facilitate bootstrapping, the root nameserver also hosts an **RMI registry** (`java.rmi.registry.Registry`), and provides a **remote object** that can then be located and called by new nameservers via RMI.

Address Lookup In contrast to the recursive algorithm used for nameserver and user address registrations, the lookup of user addresses will be handled in an iterative manner (see Figure 4 for an example): for further information see Section 5.1.4.

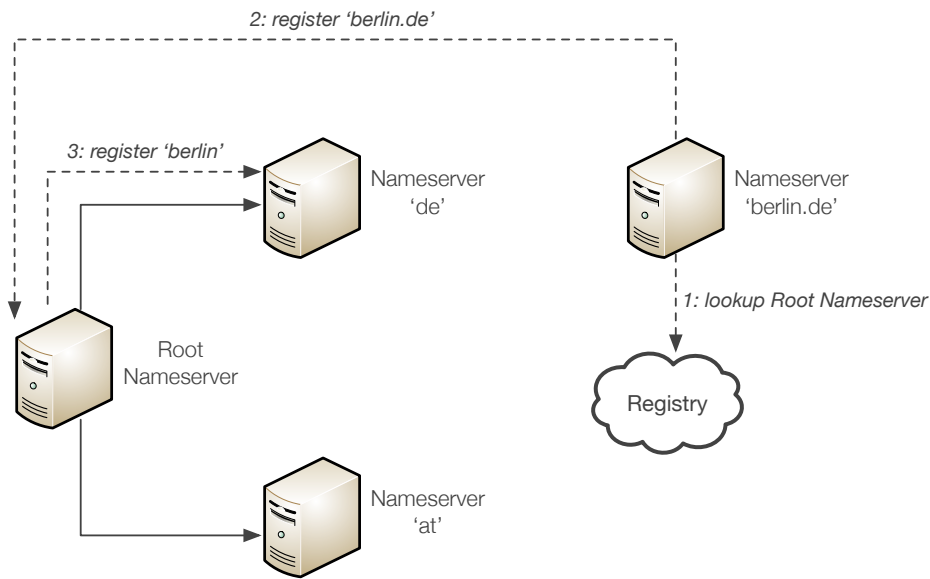


Figure 2: Naming Service: Registration

5.1.3 Nameserver

Arguments

The nameserver application reads the following parameters from the `ns-X.properties` config file, where `X` is the ID of the nameserver:

- `root_id`: the name the root nameserver is bound to or shall be bound to in case the nameserver you are currently starting is the root nameserver.
- `registry.host`: the host name or IP address where the registry is running.
- `registry.port`: the port where the RMI registry is listening for connections.
- `domain`: the domain that is managed by this nameserver. The root nameserver is the only nameserver that does not have this property. Therefore you can easily check if the nameserver you are currently starting is either an ordinary nameserver or a root nameserver.

You can assume that the parameters are valid and you do not have to verify them.

Implementation Details

Nameservers communicate in two directions, i.e., with other nameservers and the chatserver. This communication will be realized by providing remote interfaces (extensions of `java.rmi.Remote`). Methods of these interfaces can be invoked remotely over a network. The template already provides two different remote interfaces that you should implement. The interface used for the communication between nameservers and the chatserver is `INameserverForChatserver`, and the interface used between nameservers is `INameserver`¹¹.

To bind and lookup remote objects at a central location, Java RMI provides a registry service. In our case, the registry is used by the root nameserver to bind its remote object s.t. it can be located and invoked by other nameservers. The root nameserver starts a registry using `LocateRegistry.createRegistry(int port)` which creates and exports a `Registry` instance on the local host. Note that **the root nameserver is the only server that needs to create a registry**. Other nameservers can connect to the registry using `LocateRegistry.getRegistry(String host,int port)`. You should use the provided

¹¹Note that `INameserver` extends `INameserverForChatserver`, so you can combine the implementation. Having two different interfaces allows us separate concerns and limit API exposure to the chatserver.

config property `registry.port` to get the port the registry should accept requests on. Additionally, when locating the registry, use the `registry.host` property to read the host the registry is bound to.

After obtaining a reference to the `Registry`, this service can be used to bind objects that implement the `Remote` interface (using the `Registry.bind(String name, Remote obj)` method). These objects can subsequently be obtained by others using `Registry.lookup(String name)`. The RMI registry is not the only way to obtain remote objects. Remote objects can also be passed as arguments to method calls, serving as so called *callbacks*. Both approaches will be used in the implementation of our naming system.

When the root nameserver starts, it should create the registry and then bind its `INameserver` object to that registry. The config property `root_id` contains the name the object should be associated with. **The root nameserver is the only server that binds a remote object to the registry.** Both the chatserver and other nameservers can then locate this remote object via the registry, using the `root_id` config property value as name.

Instead of using the RMI registry, nameservers store reference to their child nameserver's remote objects in an appropriate in-memory structure. To register a domain, the new nameserver connects to the RMI registry, looks up the root nameserver's remote object, and provides its own remote objects as *callbacks* when invoking the `registerNameserver` method of the root nameserver's `INameserver` object. The callback objects are then passed to the next nameserver through subsequent invocations of `INameserver.registerNameserver(...)` until the correct parent nameserver is reached, as described earlier (see Figure 2).

When registering a new domain, keep in mind that the desired domain may already be in use by another server, or that an intermediary zone may not exist. In these cases, throw meaningful exceptions and pass them back to the actual requester.

As remote objects may be invoked concurrently, you have to ensure thread safety. Specifically, you should make sure that the data structures you use to manage subdomains can deal with concurrent access. You may consult the Java Concurrency Tutorial¹² to solve this problem.

You may assume that other nameservers always remain accessible, that is, you do not have to deal with zone failures. Also, data does not have to be persisted after shutting down a nameserver.

Use the nameserver's console for logging any ongoing events, e.g. what nameservers get registered, what zones are requested by the chatserver etc. An exemplary output is shown in the following:

```
17:31:13 : Registering nameserver for zone 'berlin'
17:33:45 : Nameserver for 'berlin' requested by chatserver
```

Finally, the nameserver accepts the following interactive commands:

- **!nameservers**

Prints out each known nameserver (zones) in alphabetical order, from the perspective of this nameserver.

E.g.:

```
# root nameserver
>: !nameservers
1. at
2. de

# de nameserver
>: !nameservers
1. berlin
```

- **!addresses**

Prints out some information about each stored address, containing username and address (IP:port), arranged by the username in alphabetical order. E.g.:

¹²<http://docs.oracle.com/javase/tutorial/essential/concurrency/index.html>

```
# nameserver 'at'
>: !addresses
1. maria 192.168.0.1:8888

# nameserver 'berlin.de'
>: !addresses
1. bob 192.168.0.3:7777
```

- `!exit`

Shutdown the nameserver. Do not forget to unexport its remote object using the static method `UnicastRemoteObject.unexportObject(Remote obj, boolean force)` and in the case of the root nameserver also unregister the remote object and close the registry by invoking the before mentioned static `unexportObject` method and registry reference as parameter. Otherwise the application may not stop.

5.1.4 Chatserver Updates

Arguments

There are three new configuration properties in the `chatserver.properties` file, which the chatserver will need in order to use the domain service.

- `root_id`: the name the root nameserver is bound to.
- `registry.host`: the host name or IP address where the registry is running.
- `registry.port`: the port where the RMI registry is listening for connections.

Implementation Details

At startup, the chatserver now also reads the previously mentioned additional properties to obtain the information where the RMI registry is located. Afterwards it is able to retrieve the remote reference of the root nameserver using the `root_id` property.

The classes and methods you will need for all these steps have already been explained above: `LocateRegistry.getRegistry(String host, int port)` and `Registry.lookup(String name)`.

After obtaining the root nameserver, the chatserver waits for commands sent by the client. Now to facilitate the new naming service, we have to update our approach from the first lab when handling the following commands sent by the client.

- **Register the private address of a user:** When a client wants to register the private address of a user, the chatserver now uses the distributed naming service as illustrated in Figure 3.

First, when receiving the request, the chatserver forwards it to the root nameserver by invoking the `INameserverForChatserver.registerUser(String name, String address)`. The root nameserver registers the user address in a recursive manner by forwarding the request to the next lower level in charge. This procedure continues until name resolution is no longer possible. Considering the example in Figure 3, we want to register the address for `'bob.berlin.de'`. The root nameserver first forwards the request to the `'de'` nameserver, which further forwards the request to the nameserver responsible for the `'berlin.de'` domain, which must be the requested zone. Therefore, the `'berlin.de'`-nameserver can now store the user and address. In any error case (e.g., there exists no nameserver for a requested zone) use the provided exceptions, and provide meaningful messages for the user. Furthermore, you can assume that nameserver zones and usernames will use a different naming schema, which avoids overlapping names.

- **Lookup the private address of a user:** Next, when a client wants to lookup the private address of a user, the chatserver also uses the naming service as illustrated in Figure 4.

The processing of lookup requests thus follows an iterative approach: First, the chatserver gets the remote object of the `'de'`-domain from the root nameserver (using the `INameserverForChatserver`.

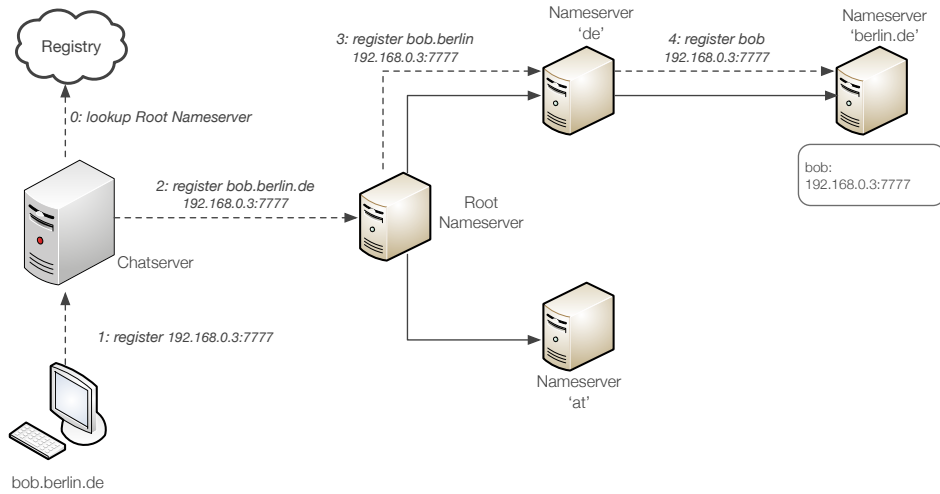


Figure 3: Naming Service: User-Address registration

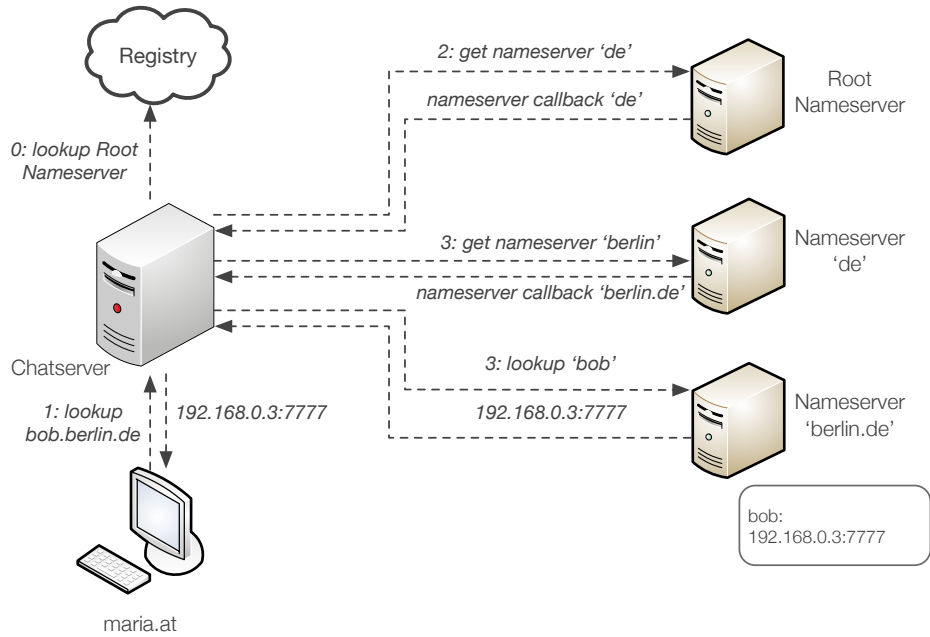


Figure 4: Naming Service: User-Address lookup

`getNameserver(String zone)`). Then the chatserver asks the 'de'-nameserver for the remote object of the 'berlin'-zone. Since this nameserver is the last subdomain in the resolution, it must be the one that handles the respective user. Therefore, the chatserver can finally ask the nameserver for the private address of the user by invoking the `INameserverForChatserver.lookup(String username)`. If a domain or the user can not be found, provide a meaningful message for the user.

5.2 Stage 2 - Secure channel (12 points)

The first stage of establishing a secure channel (for the TCP communication between the client and the chatserver) is a mutual authentication and encryption. We will authenticate a client and the chatserver using public-key cryptography. This type of authentication is explained in the book *Distributed Systems: Principles and Paradigms (2nd edition)*, page 404, Figure 9-19. However, we also describe the principles in detail below. The concrete authentication algorithm is described in the next section.

Important: Please note that in this assignment you are not allowed to use `javax.crypto.CipherInputStream` and `javax.crypto.CipherOutputStream`. Instead you should encrypt and decrypt the message communication manually using instances of `javax.crypto.Cipher`.

Your system transmits messages as plain text. However, encryption with ciphers is performed on binary data. You should therefore use **Base64 binary-to-text** encoding to convert byte arrays to strings before transmission, and vice versa upon receiving messages (see code snippets¹³). Make sure to avoid any unnecessary conversions.

We highly recommend to hide the security aspects from the rest of your application as much as possible. Note that plain sockets or encrypted channels have many commonalities, e.g., they are both used to send and receive messages. It can be beneficial to define a common communication interface that hides the details of the underlying implementation. The Decorator¹⁴ pattern can then be used to layer the encoding and encryption process. For example, you could write a `TcpChannel` that implements a common `Channel` interface and provides ways to send and receive objects over a socket. Next, you could write a `Base64Channel` class that also implements your `Channel` interface and encodes or decodes objects using Base64 before passing them to the underlying `Channel`. Following this approach may simplify your work in stage 2 and 3.

5.2.1 Authentication Algorithm

Important: We will test your submission using automated scenarios and a modified client. For this to work, you are required to implement the authentication algorithm (including the syntax of messages) **exactly** as described here. Failure to do so may result in losing points!

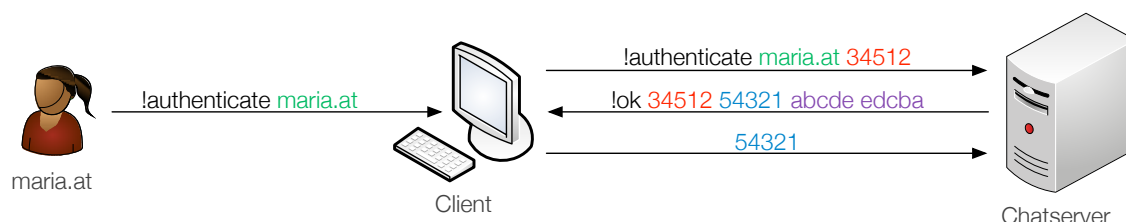


Figure 5: Secure Channel: Overview

The authentication procedure consists of three steps (see Figure 5):

1st Message: The first message is an `!authenticate` request (analogous to the `!login` command from Lab 1). The syntax of the message is: `!authenticate <username> <client-challenge>`. This message is sent by the client and is encrypted using RSA initialized with the chatserver's public key.

- The client-challenge is a 32 byte random number, which the client generates freshly for each request (see code snippets¹⁵ to learn how to generate a secure random number). Encode the challenge separately using Base64 before appending it and encrypting the overall message.
- Initialize the RSA cipher with the "RSA/NONE/OAEPWithSHA256AndMGF1Padding" algorithm.

¹³<https://tuwel.tuwien.ac.at/mod/book/view.php?id=277009&chapterid=118>

¹⁴http://en.wikipedia.org/wiki/Decorator_pattern

¹⁵<https://tuwel.tuwien.ac.at/mod/book/view.php?id=277009&chapterid=118>

- As stated above, do not forget to encode your overall encrypted message using Base64 before sending it to the chatserver.

2nd Message: The second message is sent by the chatserver and is encrypted using RSA initialized with the user's public key. The syntax is: `!ok <client-challenge> <chatserver-challenge> <secret-key> <iv-parameter>`.

- The client-challenge is the challenge that client sent in the first message. This proves to the client that the chatserver successfully decrypted the first message (i.e., it proves the chatserver's identity).
- The chatserver-challenge is also a 32 byte random number generated freshly for each request by the chatserver.
- The last two arguments are our session key. The first part is a random 256 bit secret key and the second is a random 16 byte initialization vector (IV) parameter.
- **Every argument** has to be encoded using Base64 before encrypting the overall message!
- The overall encrypted message is then encoded using Base64 and sent to the client.

When the client receives this message, it checks if the received `<client-challenge>` matches the sent one. In case the challenge does not match, the client prints an error message for the user and stops the handshake.

3rd Message: The third message is just the chatserver-challenge from the second message. This proves to the chatserver that the client successfully decrypted the second message (i.e., it further proves the client's identity). This message is the first message sent using AES encryption.

- Initialize the AES cipher using the `<secret-key>` and the `<iv-parameter>` from the second message. Details about these parameters are out of scope of this lab, but you can learn about them in a cryptography lecture.
- Use the "AES/CTR/NoPadding" algorithm for the AES cipher.
- Again, encrypt and encode the message before sending it.

When the chatserver receives this message, it checks if the received `<server-challenge>` matches the sent one. In case the challenge does not match, the chatserver prints an error message for the user and closes the connection.

The final result of the authentication procedure is an AES-encrypted secure channel between the client and the chatserver, used to encrypt all future communication. You **may not** send any message unencrypted (between a client and the chatserver) in this assignment. You **may not** send any messages besides the first two authentication messages (as described above) using the RSA encryption (i.e., the RSA encryption is strictly for the authentication part).

To generate a random 32 byte number to be used for challenges, and random 16 bytes numbers to be used for IV parameter, you should use the `java.security.SecureRandom` class and its `nextBytes()` method as shown in the Hints & Tricky Parts section¹⁵. Base64 encoding them is required because this method could return bytes which are unsuited to be inserted in a text message. The same holds true for random secret keys in the AES algorithm (see the Hints & Tricky Parts section¹⁵ on how to generate them).

Important: Always encode your challenges, IV parameters and secret keys separately in your message using Base64. The message is then encrypted and gets Base64-encoded again before sending it.

5.2.2 Client Application Behavior

When a client application is started, it doesn't know which user will try to log in. However, each user has their own public and private key. Therefore the client application will need to process the `!authenticate` request before it is sent to the chatserver, to find out which user is trying to log in and read the respective private key (used for decrypting the `!ok` message). Make sure a private key for this user exists, otherwise, print an error message. The processing of the `!authenticate` command also

includes appending the client-challenge. When implementing this stage, the `!login` command from the previous lab becomes obsolete.

There are two new configuration properties in the `client.properties` file, which the client application will need for the authentication phase:

- the `keys.dir` property denotes the directory where to look for the user's private key (named `<username>.pem`),
- and the `chatserver.key` property defines the file from where to read the public key of the chatserver.

5.2.3 Chatserver Application Behavior

The chatserver application should read its private key during the startup time. The user's public keys are read when the chatserver receives an authentication request. The user is said to be online when the authentication phase is successfully completed.

There are also two configuration properties in the `chatserver.properties` file, which the chatserver will need for the authentication phase:

- the `keys.dir` property denotes a directory where to look for user's public keys (named `<username>.pub.pem`),
- and the `key` property telling where to read the private key of the chatserver.

5.3 Stage 3 - Message Integrity (7 points)

Note: You should implement the syntax for private messages exactly as described here, for the same reasons as discussed above.

In this part of the assignment we will add an integrity check for TCP messages exchanged between the clients. However, the communication won't get encrypted - the implementation will only make sure that a third party cannot tamper with a message unnoticed. To do this, it relies on Message Authentication Codes (MACs).

Whenever a client sends a TCP request to another client and whenever a client responds, the application needs to compute a hash MAC (HMAC). To generate such a HMAC you should use SHA256 hashing ("HmacSHA256") initialized with a secret key shared between the clients in the network. See the Hints & Tricky Parts¹⁶ on how to read the shared secret key or create and initialize HMACs. After the HMAC is generated, it should be encoded using Base64. Prepend the original message with the HMAC (e.g., `<HMAC> !msg <message>`).

To verify the integrity of the message, the receiver generates a new HMAC of the received plaintext to compare it with the received one. In case of a mismatch, the behavior of the receiving client should be as follows: the respective message is printed to the standard output and the sending client is informed about the tampering (sending `<HMAC> !tampered <message>`), using the same channel the message was received. When the sending client receives this report or notices a message from a receiving client itself was changed, the client prints the incident to standard output and notifies the user accordingly.

In order to read the secret key the `client.properties` define a `hmac.key` property, which denotes from where to read the secret key.

¹⁶<https://tuwel.tuwien.ac.at/mod/book/view.php?id=277009&chapterid=118>

6 Lab Port Policy

See Lab Port Policy¹⁷ for Lab 1.

7 Regular expressions

We provide some regular expressions you can use to verify that the messages you exchange between the three applications are well-formed. This is important because we will test your program against own code. We recommend to use Java assertions for this. The provided build file enables them automatically.

```
final String B64 = "a-zA-Z0-9/+";

// --- stage II ---
// Note that the verified messages still need to be encrypted (using RSA or AES,
// respectively) and encoded using Base64!!!

// authenticate request send from the client to the chatserver
String firstMessage = ...
assert firstMessage.matches("!authenticate_\\w\\.]+_["+B64+"]{43}=") : "1st_message";

// the chatserver's response to the client
String secondMessage = ...
assert secondMessage.matches("!ok_["+B64+"]{43}=_["+B64+"]{43}=_["+B64+"]{43}=_["+B64+"]{22}=") : "2nd_message";

// the last message send by the client
String thirdMessage = ...
assert thirdMessage.matches("_["+B64+"]{43}=") : "3rd_message";

// --- stage III ---
// Private messages being exchanged between clients before the final Base64 encoding
String hashedMessage = ...
assert hashedMessage.matches("_["+B64+"]{43}=_[\\s[^\\s]]+");
```

¹⁷<https://tuwiel.tuwien.ac.at/mod/book/view.php?id=277009&chapterid=115>

8 Further Reading Suggestions

- APIs:
 - IO: IO Package API¹⁸
 - Concurrency: Thread API¹⁹, Runnable API²⁰, ExecutorService API²¹, Executors API²²
 - Java TCP Sockets: ServerSocket API²³, Socket API²⁴
 - Java Datagrams: DatagramSocket API²⁵, DatagramPacket API²⁶
- Tutorials:
 - JavaInsel Sockets Tutorial - Section 21.6²⁷, 21.7²⁸: German tutorial for using TCP sockets.
 - JavaInsel Datagrams Tutorial - Section 11.11²⁹: German tutorial for using datagram sockets (note that this chapter is from edition 7 because it has been removed in newer versions).

¹⁸<http://java.sun.com/javase/7/docs/api/index.html?java/io/package-summary.html>

¹⁹<http://java.sun.com/javase/7/docs/api/index.html?java/lang/Thread.html>

²⁰<http://java.sun.com/javase/7/docs/api/index.html?java/lang/Runnable.html>

²¹<http://java.sun.com/javase/7/docs/api/index.html?java/util/concurrent/ExecutorService.html>

²²<http://java.sun.com/javase/7/docs/api/index.html?java/util/concurrent/Executors.html>

²³<http://java.sun.com/javase/7/docs/api/index.html?java/net/ServerSocket.html>

²⁴<http://java.sun.com/javase/7/docs/api/index.html?java/net/Socket.html>

²⁵<http://java.sun.com/javase/7/docs/api/index.html?java/net/DatagramSocket.html>

²⁶<http://java.sun.com/javase/7/docs/api/index.html?java/net/DatagramPacket.html>

²⁷http://openbook.galileocomputing.de/javainsel9/javainsel_21_006.htm#mjcd64e398cec5737d9a288a4b4df04e2b

²⁸http://openbook.galileocomputing.de/javainsel9/javainsel_21_007.htm#mj1ba27dc5fdf53f527163767f188e1d2e

²⁹http://openbook.galileocomputing.de/java7/1507_11_011.html#dodtp497f87ed-dd23-48d4-80c7-7e11b3ec99d6